

Predicting Stroke Risk: Exploring the Relationship Between Demographic Factors and Health Indicators

June 30, 2023

Suhun Kim (V00952629), Min Kim (V00938382), Soyun Lee (V00904763)

Department of Computer Science, University of Victoria

1 Motivation

In modern society, where factors such as an aging population, increased isolation, and sedentary lifestyles are prevalent, the availability of immediate assistance during a stroke event may be limited. Elderly individuals living alone without a support system or the decreased social interaction due to the pandemic may contribute to delayed response times in stroke emergencies. Additionally, family medical history plays a significant role in stroke risk assessment, as genetic predispositions can increase the likelihood of stroke occurrence. Recognizing the critical importance of time in stroke, the concept of the “golden hour” highlights the significance of prompt intervention for better outcomes. Being able to anticipate the probability of a stroke in advance would prove invaluable in preventing or minimizing its impact.

Our project aims to address such challenges by developing a reliable stroke prediction model that takes into account various risk factors and individual characteristics. By integrating data on relevant parameters, our model aims to provide personalized risk assessments. This allows individuals to understand their own susceptibility to stroke and take proactive steps to reduce their risk.

```
[ ]: # import modules for ML and visualization
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: # read in data, stroke_prediction_data.csv
```

```
stroke_df = pd.read_csv('stroke_prediction_data.csv')
stroke_df.head()
```

```
[ ]:      id  gender  age  hypertension  heart_disease  ever_married  work_type  \
0    0    Male  28.0             0             0             Yes   Private
```

1	1	Male	33.0	0	0	Yes	Private
2	2	Female	42.0	0	0	Yes	Private
3	3	Male	56.0	0	0	Yes	Private
4	4	Female	24.0	0	0	No	Private

	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	Urban	79.53	31.1	never smoked	0
1	Rural	78.44	23.9	formerly smoked	0
2	Rural	103.00	40.3	Unknown	0
3	Urban	64.87	28.8	never smoked	0
4	Rural	73.36	28.8	never smoked	0

2 Data structure analysis and visualization

The data set has total 12 features. The breakdown for each columns below

id: Unique identifier for each individual.

gender: Gender of the individual (Male or Female).

age: Age of the individual.

hypertension: Whether the individual has hypertension (high blood pressure) (1 if yes, 0 if no).

heart_disease: Whether the individual has a heart disease (1 if yes, 0 if no).

ever_married: Whether the individual is married (Yes or No).

work_type: Type of work the individual is engaged in (e.g., Private, Self-employed, etc.).

Residence_type: Type of residence of the individual (Urban or Rural).

avg_glucose_level: Average glucose level in the individual's blood.

bmi: Body mass index (BMI) of the individual.

smoking_status: Smoking status of the individual (e.g., never smoked, formerly smoked, unknown).

stroke: Whether the individual has had a stroke (1 if yes, 0 if no).

Therefore, each row describe a person's medical information.

For example, the first row shows that a 28-year-old man without hypertension, heart disease or stroke and he's married, lived in Urban area, work private sector and has 79.53 average glucose level and his bmi index is 31.1.

```
[ ]: # check for missing values with info() method
```

```
stroke_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 15304 entries, 0 to 15303
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	id	15304 non-null	int64
1	gender	15304 non-null	object
2	age	15304 non-null	float64
3	hypertension	15304 non-null	int64
4	heart_disease	15304 non-null	int64

```

5  ever_married      15304 non-null  object
6  work_type        15304 non-null  object
7  Residence_type    15304 non-null  object
8  avg_glucose_level 15304 non-null  float64
9  bmi              15304 non-null  float64
10 smoking_status    15304 non-null  object
11 stroke            15304 non-null  int64
dtypes: float64(3), int64(4), object(5)
memory usage: 1.4+ MB

```

Fortunately, the stroke dataset does not contain any null values, so there is no need for an imputer to fill in missing values later on.

```

[ ]: # Generate descriptive statistics with describe() method

stroke_df.describe()

```

```

[ ]:
count    id          age  hypertension  heart_disease  \
count  15304.000000  15304.000000  15304.000000  15304.000000
mean    7651.500000   41.417708    0.049726    0.023327
std     4418.028595   21.444673    0.217384    0.150946
min         0.000000    0.080000    0.000000    0.000000
25%     3825.750000   26.000000    0.000000    0.000000
50%     7651.500000   43.000000    0.000000    0.000000
75%    11477.250000   57.000000    0.000000    0.000000
max    15303.000000   82.000000    1.000000    1.000000

count    avg_glucose_level    bmi    stroke
count    15304.000000  15304.000000  15304.000000
mean         89.039853   28.112721    0.041296
std         25.476102    6.722315    0.198981
min         55.220000   10.300000    0.000000
25%         74.900000   23.500000    0.000000
50%         85.120000   27.600000    0.000000
75%         96.980000   32.000000    0.000000
max        267.600000   80.100000    1.000000

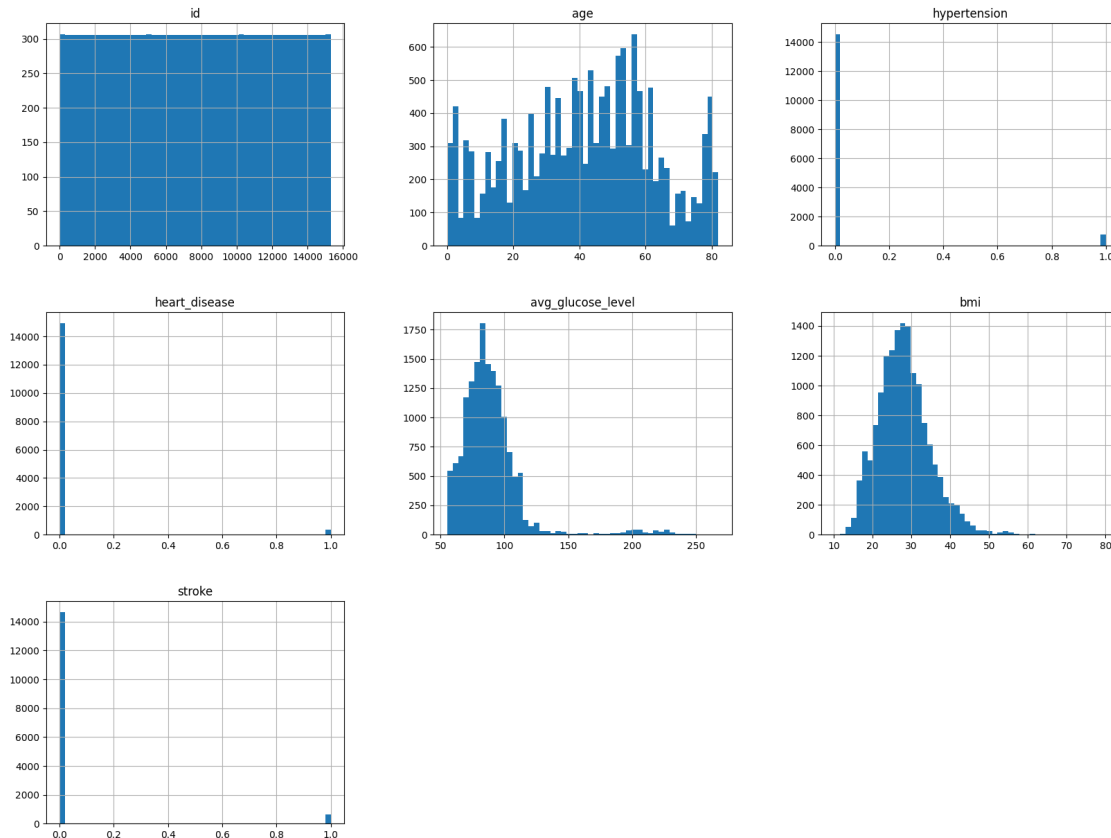
```

```

[ ]: # visualize numerical features of the data with matplotlib

stroke_df.hist(bins=50, figsize=(20,15))
plt.show()

```



After analyzing the output of the describe method and the charts, several observations can be made regarding the numerical features:

1. Based on the output of the describe method and the charts, it is clear that the BMI and average glucose level exhibit left skewness. Additionally, a significant majority of individuals have an average glucose level below 120. Therefore, these features will need transformation by using suitable methods.
2. Around 4.97% of the total population suffer from hypertension, approximately 2.33% of individuals have a recorded history of heart disease, and 4.13% of people have experienced a stroke.
3. ID column is not necessary for machine learning tasks as it's just a unique identifier, not providing meaningful data for prediction modeling.

```
[ ]: # drop id column
```

```
stroke_df.drop('id', axis=1, inplace=True)
```

```
stroke_df.head()
```

```
[ ]:   gender  age  hypertension  heart_disease  ever_married  work_type \
0    Male  28.0             0              0              0      Yes   Private
```

1	Male	33.0	0	0	Yes	Private
2	Female	42.0	0	0	Yes	Private
3	Male	56.0	0	0	Yes	Private
4	Female	24.0	0	0	No	Private

	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	Urban	79.53	31.1	never smoked	0
1	Rural	78.44	23.9	formerly smoked	0
2	Rural	103.00	40.3	Unknown	0
3	Urban	64.87	28.8	never smoked	0
4	Rural	73.36	28.8	never smoked	0

```
[ ]: # Explore categorical features of the data with value_counts() method
```

```
stroke_df.value_counts('smoking_status')
```

```
[ ]: smoking_status
never smoked      6281
Unknown           4543
formerly smoked   2337
smokes            2143
dtype: int64
```

```
[ ]: stroke_df.value_counts('work_type')
```

```
[ ]: work_type
Private           9752
children          2038
Self-employed     1939
Govt_job          1533
Never_worked       42
dtype: int64
```

```
[ ]: stroke_df.value_counts('Residence_type')
```

```
[ ]: Residence_type
Rural      7664
Urban      7640
dtype: int64
```

```
[ ]: stroke_df.value_counts('ever_married')
```

```
[ ]: ever_married
Yes      10385
No        4919
dtype: int64
```

```
[ ]: stroke_df.value_counts('gender')
```

```
[ ]: gender
      Female    9446
      Male      5857
      Other       1
      dtype: int64
```

By examining the output of the describe and value_counts methods, we can draw several observations about the categorical features.

1. The Residence_type feature is almost evenly distributed between two categories, namely Rural and Urban.
2. The ever_married feature consists of two unique values, namely Yes and No. However, the majority of individuals in this column are married.
3. The gender feature predominantly consists of females (9446) and fewer males represented in the dataset. There's only one record categorized as "Other" gender.
4. Approximately 40% of individuals responded that they have never smoked, while around 15% of individuals reported currently smoking and another 15% having smoked in the past.
5. The most prevalent work type is the private sector, with a majority of individuals engaged in this category. Additionally, there is a notable presence of approximately 2000 individuals labeled as "children," suggesting they are below the legal working age. A similar number of individuals are self-employed. A small group of only 42 individuals responded that they have never worked in their life.

Based on the observation, it is decided to drop the row corresponding to the "Other" gender individual. Including this category in the prediction model, for instance, through one-hot encoding, would introduce an additional column and potentially complicate the model. As a result, it may lead to computationally expensive processes or potentially inaccurate results.

```
[ ]: # drop other row from gender column as it's too edge case

stroke_df.drop(stroke_df[stroke_df['gender']=='Other'].index, axis=0,
               inplace=True)
```

```
[ ]: # Double check if row was dropped and new column was added

stroke_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 15303 entries, 0 to 15303
Data columns (total 11 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   gender          15303 non-null  object
 1   age             15303 non-null  float64
 2   hypertension    15303 non-null  int64
```

```

3   heart_disease      15303 non-null  int64
4   ever_married      15303 non-null  object
5   work_type         15303 non-null  object
6   Residence_type     15303 non-null  object
7   avg_glucose_level  15303 non-null  float64
8   bmi               15303 non-null  float64
9   smoking_status     15303 non-null  object
10  stroke            15303 non-null  int64
dtypes: float64(3), int64(3), object(5)
memory usage: 1.4+ MB

```

```
[ ]: stroke_df.corr()
```

```

[ ]:
          age  hypertension  heart_disease  avg_glucose_level  \
age          1.000000      0.232624      0.193394      0.108933
hypertension  0.232624      1.000000      0.074176      0.122838
heart_disease 0.193394      0.074176      1.000000      0.137151
avg_glucose_level 0.108933  0.122838      0.137151      1.000000
bmi           0.390630      0.106651      0.054571      0.108177
stroke        0.261319      0.145892      0.111516      0.149278

          bmi  stroke
age          0.390630  0.261319
hypertension 0.106651  0.145892
heart_disease 0.054571  0.111516
avg_glucose_level 0.108177  0.149278
bmi           1.000000  0.065377
stroke        0.065377  1.000000


```

```

[ ]: # Visualize correlation matrix with heatmap

plt.figure(figsize=(12,8))
sns.heatmap(stroke_df.corr(), annot=True, cmap='coolwarm')
plt.show()

```




Based on the visualized correlation matrix, it is observed that hypertension and heart_disease are significant factors contributing to strokes, except age.


However, as previously mentioned, hypertension and heart_disease are binary variables. To highlight the importance of these features and introduce more diversity, it is decided to generate an additional column indicating the presence of either no hypertension or heart disease, only one of them, or both conditions.

```
[ ]: # Extra column for hypertension and heart disease in one column.
# hypertension AND heart_disease: 2, hypertension OR heart_disease: 1, no
↳hypertension or heart_disease: 0

stroke_df['hypertension_heart_disease'] = stroke_df['hypertension'] +
↳stroke_df['heart_disease']

[ ]: sns.scatterplot(data=stroke_df, x='age', y='bmi', hue='stroke', style="stroke",
↳size="stroke", palette=sns.color_palette(['#aaf9f4', 'red']), alpha=.9,
↳sizes=(200,20), markers="+")
plt.show()
sns.scatterplot(data=stroke_df, x='age', y='avg_glucose_level', style="stroke",
↳hue='stroke', palette=sns.color_palette(['#aaf9f4', 'blue']), alpha=.9,
↳sizes=(200,20), markers="+")
plt.show()
```



The highest correlation is observed between age and stroke, and there is also a moderate positive correlation between age and BMI. By examining the chart with age on the x-axis and BMI on the y-axis (with the stroke variable indicated by hue), observations below can be made.

The stroke rate dramatically increases for individuals who are over 50 years old. Additionally, for individuals over 50, the BMI does not significantly impact the probability of stroke. However, for individuals under 50 years old, BMI has a more noticeable effect, particularly for those with a BMI over 30.

The correlation of `avg_glucose_level` with other features appears to be relatively consistent. However, when plotted against age and stroke, it becomes evident that `avg_glucose_level` is distinctly separated into two groups: one below 125 and another above 180. Notably, individuals with glucose levels exceeding 180 are likely to be diagnosed with diabetes. Moreover, the prevalence of diabetes significantly increases in individuals over 50 years old. Therefore, it can be concluded that individuals who are both over 50 years old and have diabetes are at a higher risk of experiencing a stroke.

Based on these observations, the decision was made to create two additional columns: `age^2 * bmi` and `age^2 * avg_glucose_level`


```
[ ]: # additional columns for age^2 * bmi and age^2 * avg_glucose_level if age is
      ↳ over 50
```


```
stroke_df['age_bmi'] = 0
stroke_df['age_avg_glucose_level'] = 0

stroke_df.loc[stroke_df['age'] >= 50, ['age_bmi']] = (stroke_df['age'] - 50) ** 2 * stroke_df['bmi']
stroke_df.loc[stroke_df['age'] >= 50, ['age_avg_glucose_level']] = (stroke_df['age'] - 50) ** 2 * stroke_df['avg_glucose_level']
```

```
[ ]: stroke_df[['age_bmi']].hist(bins=30, figsize=(20,10))
stroke_df[['age_avg_glucose_level']].hist(bins=50, figsize=(20,10))

plt.show()
```





3 Data preparation

The previously mentioned histograms make it clear that there is a significant correlation between BMI, average glucose level, and people who are 50 years of age or older. It can now be divided the data into train and test sets by introducing the newly created column `stroke_df`. Attribute interdependencies are examined while analyzing the trainset data. Notably, the correlation between age and average glucose level is 0.348894, while the correlation between age and BMI level is 0.329024. These results show a definite correlation between these variables.

```
[ ]: # split data into train and test sets

from sklearn.model_selection import train_test_split

stroke_train_set, stroke_test_set = train_test_split(stroke_df, test_size=0.2,
↳ random_state=42)

stroke_train_set.shape, stroke_test_set.shape
```

```
[ ]: ((12242, 14), (3061, 14))
```

```
[ ]: num_attribs = ['age', 'hypertension', 'heart_disease', 'avg_glucose_level',
↳ 'age_bmi', 'hypertension_heart_disease', 'age_avg_glucose_level', 'stroke']

stroke_train_set[num_attribs].corr()['stroke'].sort_values(ascending=False)
```

```
[ ]: stroke                1.000000
age_avg_glucose_level    0.348894
age_bmi                  0.329024
```

```

age                                0.259873
hypertension_heart_disease         0.171756
hypertension                       0.150190
avg_glucose_level                  0.145506
heart_disease                      0.095446
Name: stroke, dtype: float64

```

```
[ ]: # create a copy of the training set and test set by dropping the target column
```

```

stroke_train = stroke_train_set.drop('stroke', axis=1)
stroke_train_labels = stroke_train_set['stroke'].copy()

stroke_test = stroke_test_set.drop('stroke', axis=1)
stroke_test_labels = stroke_test_set['stroke'].copy()

# check the shape of the training set and test set

stroke_train.shape, stroke_test.shape

```

```
[ ]: ((12242, 13), (3061, 13))
```

```
[ ]: # create a imputer for numerical attributes
# create a imputer for categorical attributes
```

```

from sklearn.impute import SimpleImputer

num_imputer = SimpleImputer(strategy='median')
cat_imputer = SimpleImputer(strategy='most_frequent')

```

```
[ ]: # create a pipeline for categorical attributes
```

```

from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import OrdinalEncoder
from sklearn.pipeline import Pipeline

cat_attribs = ['gender', 'ever_married', 'work_type', 'Residence_type',
               ↪ 'smoking_status']

cat_encoder = OrdinalEncoder()

cat_pipeline = Pipeline([('cat_encoder', cat_encoder)])

cat_pipeline

```

```
[ ]: Pipeline(steps=[('cat_encoder', OrdinalEncoder())])
```

```
[ ]: # scale numerical attributes

from sklearn.preprocessing import StandardScaler

num_attribs = ['age_bmi', 'hypertension_heart_disease', 'age_avg_glucose_level']


# create pipeline for numerical attributes


std_scaler = Pipeline([('std_scaler', StandardScaler())])

std_age_avg_glucose_level = std_scaler.
    ↪fit_transform(stroke_train[['age_avg_glucose_level']])
std_age_bmi = std_scaler.fit_transform(stroke_train[['age_bmi']])

plot = sns.histplot(std_age_avg_glucose_level)
plt.show()

plot = sns.histplot(std_age_bmi)
plt.show()
```





```
[ ]: # transform avg_glucose_level and bmi columns with log transformation



from sklearn.preprocessing import FunctionTransformer

log_transformer = FunctionTransformer(np.log, inverse_func=np.exp,
    ↪feature_names_out="one-to-one")

log_avg_glucose_level = log_transformer.
    ↪fit_transform(stroke_train[['avg_glucose_level']])
log_bmi = log_transformer.fit_transform(stroke_train[['bmi']])

plot = sns.histplot(log_avg_glucose_level)
plt.show()

plot = sns.histplot(log_bmi)
plt.show()
```




```
[ ]: # use minmax scaler to scale numerical attributes

from sklearn.preprocessing import MinMaxScaler

minmax_scaler = MinMaxScaler()

# make a pipeline for minmax scaler


minmax_transtormer = Pipeline([('minmax_scaler', minmax_scaler)])

log_avg_glucose = minmax_transtormer.
    ↳fit_transform(stroke_train[["avg_glucose_level"]])

plot = sns.histplot(log_avg_glucose)

plot.set(xlabel='log(avg_glucose_level)', ylabel='count')

plt.show()
```



```
[ ]: # use power transformation to transform avg_glucose_level

from sklearn.preprocessing import PowerTransformer


power_transformer = PowerTransformer()

power_avg_glucose_level = power_transformer.
    ↪fit_transform(stroke_train[['avg_glucose_level']])

plot = sns.histplot(power_avg_glucose_level)

plot.set(xlabel='avg_glucose_level', ylabel='count')

plt.show()
```



In order to normalize the bmi and avg_glucose_level columns, various methods were applied, such as StandardScaler, logarithm function, MinMaxScaler, and PowerTransformer. The logarithm function was successful in achieving a more balanced distribution for the bmi column, but it did not significantly improve the distribution of the avg_glucose_level column. Similarly, using MinMaxScaler did not yield the desired results.

Finally, PowerTransformer was found to be effective in scaling the data and reducing the tail length

of the distributions. However, since it is important to preserve the tail of the distributions in the age_avg_glucose_level and age_bmi columns to maintain their specific characteristics and original meanings, the StandardScaler will be used instead. This approach will ensure that the length of the tail is retained while effectively normalizing the data.

```
[ ]: # create a full pipeline for numerical and categorical attributes

from sklearn.compose import ColumnTransformer

num_std_trans_pipeline = Pipeline([('num_imputer', num_imputer),
                                   ('std_scaler', std_scaler)])

num_log_trans_pipeline = Pipeline([('num_imputer', num_imputer),
                                   ('log_transformer', log_transformer)])

num_power_trans_pipeline = Pipeline([('num_imputer', num_imputer),
                                     ('power_transformer', power_transformer)])

num_minmax_trans_pipeline = Pipeline([('num_imputer', num_imputer),
                                      ('minmax_transtormer',
                                       ↪minmax_transtormer)])

num_pipeline = Pipeline([('num_imputer', num_imputer)])

cat_pipeline = Pipeline([('cat_imputer', cat_imputer),
                         ('cat_encoder', cat_encoder)])

preprocessing = ColumnTransformer(transformers=[('num_std',
↪num_std_trans_pipeline, ['age']),
                                              ('num_log',
↪num_log_trans_pipeline, ['bmi']),
                                              ('num_power',
↪num_power_trans_pipeline, ['avg_glucose_level', 'age_avg_glucose_level']),
                                              ('num_minmax',
↪num_minmax_trans_pipeline, ['age_bmi', 'age_avg_glucose_level']),
                                              ('num', num_pipeline,
↪['hypertension', 'heart_disease', 'hypertension_heart_disease']),
                                              ('cat', cat_pipeline,
↪cat_attribs)])

preprocessing
```

```
[ ]: ColumnTransformer(transformers=[('num_std',
                                   Pipeline(steps=[('num_imputer',
SimpleImputer(strategy='median')),
                                   ('std_scaler',
                                   Pipeline(steps=[('std_scaler',
```

```

StandardScaler()))]])),
                                ['age']],
                                ('num_log',
                                 Pipeline(steps=[('num_imputer',
SimpleImputer(strategy='median')),
                                ('log_transformer',
FunctionTransformer(feature_names_out='one-to-one',
func=<ufunc 'log'>,
                                in...
                                ['age_bmi', 'age_avg_glucose_level']],
                                ('num',
                                 Pipeline(steps=[('num_imputer',
SimpleImputer(strategy='median'))])),
                                ['hypertension', 'heart_disease',
                                'hypertension_heart_disease']],
                                ('cat',
                                 Pipeline(steps=[('cat_imputer',
SimpleImputer(strategy='most_frequent')),
                                ('cat_encoder',
                                 OrdinalEncoder()))]),
                                ['gender', 'ever_married', 'work_type',
                                'Residence_type', 'smoking_status']]])

```

3.0.1 Train model and predict on test set

```

[ ]: # train a logistic regression model

from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import make_pipeline

log_reg = make_pipeline(preprocessing, LogisticRegression(random_state=42,
    ↳max_iter=1000))

log_reg.fit(stroke_train, stroke_train_labels)

predictions = log_reg.predict(stroke_train)

# evaluate the model including confusion matrix

from sklearn.metrics import confusion_matrix, precision_score, recall_score,
    ↳f1_score, roc_auc_score, roc_curve

print(confusion_matrix(stroke_train_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_train_labels, predictions,
    ↳labels=[True, False]))

```

```

print('recall_score: ', recall_score(stroke_train_labels, predictions,
    ↳labels=[True, False]))

# evaluate on the test set

predictions = log_reg.predict(stroke_test)

print(confusion_matrix(stroke_test_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_test_labels, predictions,
    ↳labels=[True, False]))

print('recall_score: ', recall_score(stroke_test_labels, predictions,
    ↳labels=[True, False]))

```

```

[[ 10  488]
 [ 11 11733]]
precision score:  0.47619047619047616
recall_score:  0.020080321285140562
[[  3  131]
 [  5 2922]]
precision score:  0.375
recall_score:  0.022388059701492536

```

```

[ ]: # train a random forest model

from sklearn.ensemble import RandomForestClassifier

forest_clf = make_pipeline(preprocessing,
    ↳RandomForestClassifier(random_state=42))

forest_clf.fit(stroke_train, stroke_train_labels)

predictions = forest_clf.predict(stroke_train)

print(confusion_matrix(stroke_train_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_train_labels, predictions,
    ↳labels=[True, False]))

print('recall_score: ', recall_score(stroke_train_labels, predictions,
    ↳labels=[True, False]))

# evaluate on the test set

predictions = forest_clf.predict(stroke_test)

```

```

print(confusion_matrix(stroke_test_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_test_labels, predictions,
↪labels=[True, False]))

print('recall_score: ', recall_score(stroke_test_labels, predictions,
↪labels=[True, False]))

```

```

[[ 496    2]
 [    0 11744]]
precision score:  1.0
recall_score:  0.9959839357429718
[[    6   128]
 [   15 2912]]
precision score:  0.2857142857142857
recall_score:  0.04477611940298507

```

```

[ ]: # train decision tree model

from sklearn.tree import DecisionTreeClassifier

tree_clf = make_pipeline(preprocessing, DecisionTreeClassifier(random_state=42))

tree_clf.fit(stroke_train, stroke_train_labels)

predictions = tree_clf.predict(stroke_train)

print(confusion_matrix(stroke_train_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_train_labels, predictions,
↪labels=[True, False]))

print('recall_score: ', recall_score(stroke_train_labels, predictions,
↪labels=[True, False]))

# evaluate on the test set

predictions = tree_clf.predict(stroke_test)

print(confusion_matrix(stroke_test_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_test_labels, predictions,
↪labels=[True, False]))

print('recall_score: ', recall_score(stroke_test_labels, predictions,
↪labels=[True, False]))

```

```

[[ 498    0]

```

```

[ 0 11744]]
precision score: 1.0
recall_score: 1.0
[[ 21 113]
 [111 2816]]
precision score: 0.1590909090909091
recall_score: 0.15671641791044777

```

Three different training models were applied to the dataset, and the pipeline of each attribute in the trainset was connected. The decision regarding the train models was made in light of the information learned from the SENG474 lecture. The Logistic Regression train model, which combined the Logistic Regression technique with a confusion matrix, was the first one used. The Decision Tree and Random Forest Classifier models used the same methodology.

Since accuracy is primarily used to evaluate the performance of machine learning datasets in the majority of cases discussed during classes and lab sessions, the confusion matrix and visualising the precision and recall scores over accuracy was chosen. However, accuracy is not the only important parameter in the context of our topic, which deals with heart stroke. The main goal is to reduce the confusion matrix's false negative value. Accepting false positives can offer cautionary information about people's health conditions, but doing so risks undermining the project's primary goal of identifying those at risk of suffering a heart stroke. As a result, recall is significant, and the confusion matrix is a useful tool for examining the predictions connected to each value.

Both precision and recall scores for the train dataset and test dataset were low, according to the findings of the Logistic Regression model. On the trainset, the Random Forest Classifier model achieved 100% precision and 99.5984% recall, but on the testset, it only managed 28.5% precision and 4.4776% recall, demonstrating overfitting. Finally, the Decision Tree Classifier model produced the highest recall rate so far, 15.9% and 15.6716%, while overfitting the testset with 100% precision.

```

[ ]: # train grid search cv on random forest model

from sklearn.model_selection import GridSearchCV

param_grid = [{'randomforestclassifier__n_estimators': [3, 10, 30],
↳ 'randomforestclassifier__max_features': [2, 4, 6, 8]},
               {'randomforestclassifier__bootstrap': [False],
↳ 'randomforestclassifier__n_estimators': [3, 10],
↳ 'randomforestclassifier__max_features': [2, 3, 4]]

grid_search = GridSearchCV(forest_clf, param_grid, cv=5, scoring='roc_auc',
↳ return_train_score=True)

grid_search.fit(stroke_train, stroke_train_labels)

predictions = grid_search.predict(stroke_train)

print(confusion_matrix(stroke_train_labels, predictions, labels=[True, False]))

```

```

print('precision score: ', precision_score(stroke_train_labels, predictions,
↳labels=[True, False]))

print('recall_score: ', recall_score(stroke_train_labels, predictions,
↳labels=[True, False]))

# evaluate on the test set

predictions = grid_search.predict(stroke_test)

print(confusion_matrix(stroke_test_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_test_labels, predictions,
↳labels=[True, False]))

print('recall_score: ', recall_score(stroke_test_labels, predictions,
↳labels=[True, False]))

```

```

[[ 478   20]
 [   0 11744]]
precision score:  1.0
recall_score:  0.9598393574297188
[[   8  126]
 [  25 2902]]
precision score:  0.24242424242424243
recall_score:  0.05970149253731343

```

```

[ ]: # train random search cv on random forest model

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

param_distributions = {'randomforestclassifier__n_estimators': randint(low=1,
↳high=200),
                        'randomforestclassifier__max_features': randint(low=1,
↳high=8)}

rnd_search = RandomizedSearchCV(forest_clf, param_distributions=param_distributions,
↳n_iter=10, cv=5, scoring='roc_auc', random_state=42)

rnd_search.fit(stroke_train, stroke_train_labels)

predictions = rnd_search.predict(stroke_train)

print(confusion_matrix(stroke_train_labels, predictions, labels=[True, False]))

```



```

print('precision score: ', precision_score(stroke_train_labels, predictions,
↪labels=[True, False]))

print('recall_score: ', recall_score(stroke_train_labels, predictions,
↪labels=[True, False]))

# evaluate on the test set

predictions = rnd_search.predict(stroke_test)

print(confusion_matrix(stroke_test_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_test_labels, predictions,
↪labels=[True, False]))

print('recall_score: ', recall_score(stroke_test_labels, predictions,
↪labels=[True, False]))

```

```

[[ 497    1]
 [    0 11744]]
precision score:  1.0
recall_score:  0.9979919678714859
[[    8   126]
 [   13 2914]]
precision score:  0.38095238095238093
recall_score:  0.05970149253731343

```

Additional models, including GridSearchCV and RandomizedSearchCV, were tested in an effort to improve the results for the fine-tuning procedure. These models were chosen because they can increase generalisation, optimise model performance, and boost the accuracy of predictions. Particularly, RandomizedSearchCV offered GridSearchCV an effective and adaptable substitute.

However, it was found that both models on the train dataset showed overfitting, resulting in a recall score of 5.9701% on the test dataset.

4 Additional observation with new external factors: medical information of stroke

To enhance the predictive power of the model, new features were created or existing features were given additional weight based on medical experiment results.

1. A >11-year age difference in stroke incidence was found between smokers and nonsmokers, with a median age of 60.2 years for current smokers and 71.6 years for nonsmokers.^[1]
=> weight + 11.4 for 'smokes'. and few number(?) for formerly smoked, calculate median number for weighting number (and add or ignore for unknown)
 - never smoked 6281
 - Unknown 4543
 - formerly smoked 2337

- smokes 2143
2. As we have average glucose level of people, mean value of Fasting Blood Sugar Test and Glucose Tolerance Test used to determine if a person has diebete or not.^[2]
(i.e. a person is normal if $\text{avg_glucose_level} < 120$, prediabetes if $120 < \text{avg_glucose_level} < 162.5$, or diabetes if $\text{avg_glucose_level} > 162.5$)
Plus, according to American Diabetes Association, chances of having a stroke are 2 times higher if a person is suffering from diabetes.^[3]
=> Group people by avg_glucose_level as normal, prediabetes, diabetes. Change column as categorical attributes
 3. BMIs of 30 or above were associated with an adjusted relative risk of 2.00 (95% confidence interval [CI], 1.48-2.71) for overall stroke compared to persons with BMIs of less than 23. ^[4]
=> Therefore, rather than using a numerical value, we can divide into three distinct groups. The first group is made up of people with BMIs under 23, followed by those with BMIs between 23 and 29, and those with BMIs beyond 30.

```
[ ]: # Add 11.4 to age column if a smoking status is 'smokes' or 'formerly smoked'

stroke_train.loc[stroke_train['smoking_status'] == 'smokes', 'age'] += 11.4
stroke_train.loc[stroke_train['smoking_status'] == 'formerly smokes', 'age'] += 11.4 / 2

stroke_test.loc[stroke_test['smoking_status'] == 'smokes', 'age'] += 11.4
stroke_test.loc[stroke_test['smoking_status'] == 'formerly smokes', 'age'] += 11.4 / 2

[ ]: # if avg_glucose_level < 120, prediabetes if 120 <= avg_glucose_level < 162.5,
    or diabetes if avg_glucose_level > 162.5
# Group people by avg_glucose_level as normal, prediabetes, diabetes. Change
    column as categorical attributes

stroke_train.loc[stroke_train['avg_glucose_level'] < 120,
    'avg_glucose_level_str'] = 'normal'
stroke_train.loc[(stroke_train['avg_glucose_level'] >= 120) &
    (stroke_train['avg_glucose_level'] < 162.5), 'avg_glucose_level_str'] =
    'prediabetes'
stroke_train.loc[stroke_train['avg_glucose_level'] >= 162.5,
    'avg_glucose_level_str'] = 'diabetes'

stroke_test.loc[stroke_test['avg_glucose_level'] < 120,
    'avg_glucose_level_str'] = 'normal'
stroke_test.loc[(stroke_test['avg_glucose_level'] >= 120) &
    (stroke_test['avg_glucose_level'] < 162.5), 'avg_glucose_level_str'] =
    'prediabetes'
stroke_test.loc[stroke_test['avg_glucose_level'] >= 162.5,
    'avg_glucose_level_str'] = 'diabetes'
```

```
[ ]: # if bmi < 23, normal, if 23 <= bmi < 30, overweight, if bmi >= 30, obese
```

```
stroke_train.loc[stroke_train['bmi'] < 23, 'bmi_str'] = 'normal'
stroke_train.loc[(stroke_train['bmi'] >= 23) & (stroke_train['bmi'] < 30),
↳ 'bmi_str'] = 'overweight'
stroke_train.loc[stroke_train['bmi'] >= 30, 'bmi_str'] = 'obese'

stroke_test.loc[stroke_test['bmi'] < 23, 'bmi_str'] = 'normal'
stroke_test.loc[(stroke_test['bmi'] >= 23) & (stroke_test['bmi'] < 30),
↳ 'bmi_str'] = 'overweight'
stroke_test.loc[stroke_test['bmi'] >= 30, 'bmi_str'] = 'obese'
```

```
[ ]: stroke_train.head(10)
```

```
[ ]:      gender  age  hypertension  heart_disease  ever_married  work_type \
8304    Male  45.0             0             0             Yes  Govt_job
10732   Male  62.0             0             0             Yes   Private
7015   Female  11.0             0             0             No   children
12543   Male  56.0             0             0             Yes  Govt_job
2455    Male   8.0             0             0             No   children
1595   Female  46.0             0             0             Yes   Private
10218   Female  35.0             0             0             Yes   Private
13294   Female  22.0             0             0             No   Private
8278    Male  38.0             0             0             Yes   Private
9939   Female  63.0             0             0             Yes  Self-employed
```

```
      Residence_type  avg_glucose_level  bmi  smoking_status \
8304             Urban             92.23  32.0           Unknown
10732            Rural             93.88  43.3           Unknown
7015            Urban             88.19  15.5  never smoked
12543            Urban             88.81  33.5  never smoked
2455            Urban             90.22  18.8           Unknown
1595            Urban            111.94  25.6  never smoked
10218            Urban             77.12  24.5  never smoked
13294            Rural            102.05  21.7  never smoked
8278            Rural             73.78  24.1           Unknown
9939            Rural             99.96  28.4  never smoked
```

```
      hypertension_heart_disease  age_bmi  age_avg_glucose_level \
8304                        0      0.0              0.00
10732                       0    6235.2            13518.72
7015                        0      0.0              0.00
12543                       0    1206.0            3197.16
2455                        0      0.0              0.00
1595                        0      0.0              0.00
10218                       0      0.0              0.00
13294                       0      0.0              0.00
```

8278	0	0.0	0.00
9939	0	4799.6	16893.24

	avg_glucose_level_str	bmi_str
8304	normal	obese
10732	normal	obese
7015	normal	normal
12543	normal	obese
2455	normal	normal
1595	normal	overweight
10218	normal	overweight
13294	normal	normal
8278	normal	overweight
9939	normal	overweight

```
[ ]: # drop smoking_status, avg_glucose_level, bmi

stroke_train.drop(['smoking_status', 'avg_glucose_level', 'bmi'], axis=1,
                  inplace=True)
stroke_test.drop(['smoking_status', 'avg_glucose_level', 'bmi'], axis=1,
                 inplace=True)

stroke_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 12242 entries, 8304 to 7270
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   gender                                12242 non-null  object
1   age                                   12242 non-null  float64
2   hypertension                          12242 non-null  int64
3   heart_disease                        12242 non-null  int64
4   ever_married                         12242 non-null  object
5   work_type                            12242 non-null  object
6   Residence_type                       12242 non-null  object
7   hypertension_heart_disease           12242 non-null  int64
8   age_bmi                              12242 non-null  float64
9   age_avg_glucose_level                12242 non-null  float64
10  avg_glucose_level_str                12242 non-null  object
11  bmi_str                              12242 non-null  object
dtypes: float64(3), int64(3), object(6)
memory usage: 1.2+ MB
```

```
[ ]: # create new preprocessing pipeline
```

```

cat_attribs = ['gender', 'ever_married', 'work_type', 'Residence_type',
               ↪ 'avg_glucose_level_str', 'bmi_str']

new_reprocessing = ColumnTransformer(transformers=[('age',
               ↪ num_std_trans_pipeline, ['age']),
               ('num', num_pipeline,
               ↪ ['hypertension', 'heart_disease', 'hypertension_heart_disease']),
               ('cat', cat_pipeline,
               ↪ cat_attribs)])

stroke_train_prepared = new_reprocessing.fit_transform(stroke_train)

df_stroke_train_prepared = pd.DataFrame(stroke_train_prepared,
               ↪ columns=new_reprocessing.get_feature_names_out())

df_stroke_train_prepared.head()

```

```

[ ]:   age__age  num__hypertension  num__heart_disease  \
0  0.091335          0.0          0.0
1  0.860056          0.0          0.0
2 -1.446107          0.0          0.0
3  0.588743          0.0          0.0
4 -1.581764          0.0          0.0

      num__hypertension_heart_disease  cat__gender  cat__ever_married  \
0                0.0          1.0          1.0
1                0.0          1.0          1.0
2                0.0          0.0          0.0
3                0.0          1.0          1.0
4                0.0          1.0          0.0

      cat__work_type  cat__Residence_type  cat__avg_glucose_level_str  \
0                0.0          1.0          1.0
1                2.0          0.0          1.0
2                4.0          1.0          1.0
3                0.0          1.0          1.0
4                4.0          1.0          1.0

      cat__bmi_str
0                1.0
1                1.0
2                0.0
3                1.0
4                0.0

```

```
[ ]: # train a random forest model

from sklearn.ensemble import RandomForestClassifier

forest_clf = make_pipeline(new_reprocessing,
    ↳RandomForestClassifier(random_state=42))

forest_clf.fit(stroke_train, stroke_train_labels)

predictions = forest_clf.predict(stroke_train)

print(confusion_matrix(stroke_train_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_train_labels, predictions,
    ↳labels=[True, False]))

print('recall_score: ', recall_score(stroke_train_labels, predictions,
    ↳labels=[True, False]))

# evaluate on the test set

predictions = forest_clf.predict(stroke_test)

print(confusion_matrix(stroke_test_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_test_labels, predictions,
    ↳labels=[True, False]))

print('recall_score: ', recall_score(stroke_test_labels, predictions,
    ↳labels=[True, False]))
```

```
[[ 270  228]
 [  23 11721]]
precision score:  0.9215017064846417
recall_score:  0.5421686746987951
[[ 15  119]
 [  43 2884]]
precision score:  0.25862068965517243
recall_score:  0.11194029850746269
```

```
[ ]: # train decision tree model

from sklearn.tree import DecisionTreeClassifier

tree_clf = make_pipeline(new_reprocessing,
    ↳DecisionTreeClassifier(random_state=42))
```

```

tree_clf.fit(stroke_train, stroke_train_labels)

predictions = tree_clf.predict(stroke_train)

print(confusion_matrix(stroke_train_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_train_labels, predictions,
↪labels=[True, False]))

print('recall_score: ', recall_score(stroke_train_labels, predictions,
↪labels=[True, False]))

# evaluate on the test set

predictions = tree_clf.predict(stroke_test)

print(confusion_matrix(stroke_test_labels, predictions, labels=[True, False]))

print('precision score: ', precision_score(stroke_test_labels, predictions,
↪labels=[True, False]))

print('recall_score: ', recall_score(stroke_test_labels, predictions,
↪labels=[True, False]))

```

```

[[ 257  241]
 [  10 11734]]
precision score:  0.9625468164794008
recall_score:  0.5160642570281124
[[  18  116]
 [  67 2860]]
precision score:  0.21176470588235294
recall_score:  0.13432835820895522

```

- Comparison for random forest classifier:

The previous random forest classifier model demonstrated outstanding performance on the train set, achieving a precision score of 100% and a recall score of 99.598%. However, its performance on the test set was significantly poorer, with a precision score of 28.57% and a recall score of 4.5%. This suggests that the model overfit the training data and struggled to generalize to new, unseen data.

In contrast, a new model incorporating additional medical knowledge exhibited slightly inferior performance on the train set, particularly with a notable drop in recall score by over 40%, while experiencing only a slight decrease in precision. However, this new model showed improved performance on the test set, particularly in terms of recall, which increased from approximately 4.5% to 11.19%, while the precision score slightly decreased.

Overall, the new model demonstrated better generalization ability by achieving higher recall on the test set compared to the previous model. Although it experienced a slight decline in precision, the improved recall indicates that it successfully identified more positive cases correctly. Further model optimization and fine-tuning could potentially enhance the overall

performance of the new model on both the train and test sets.

- **Comparison for decision tree classifier:**

The previous decision tree classifier model achieved perfect performance on the train set, similar to the random forest model, with 100% precision and recall. Not surprisingly, its performance also declined dramatically on the test set. The precision and recall values on the test set are only 15.9% and 15.6%, respectively. Similar to random forest classifier model, it can be confirmed that this classifier also was overfitted on train set.

The new decision tree classifier model maintained a high precision of 96.25% and a recall of 51.6% on the train set. However, similar to the previous models, its performance dropped on the test set, with a precision of 21.18% and a recall of 13.4%.

Both classifiers demonstrated a pattern of overfitting on the train set, regardless of whether they were simple machine learning models or adjusted models with medical knowledge. Interestingly, the precision scores slightly decreased compared to the previous models, while the recall scores were significantly reduced for the train set. The precision scores for the test set did not undergo significant changes. On the contrary, the random forest classifier which doubled its recall score, while the decision tree classifier's recall score dropped.

5 Result

During the initial modeling phase, it was observed that all the classifiers achieved high accuracy on the training data but struggled to generalize to new data. This phenomenon, known as overfitting, occurs when models become too specific to the training data and fail to perform well on unseen examples. To emphasize improving the recall values, which measure the ability to correctly identify positive cases, extensive research was conducted by investigating few studies and scientific papers for potential insights. However, the complexity of the human body posed challenges, and the incorporation of new factors did not yield the significant improvements anticipated. The introduction of these additional factors resulted in a decline in recall scores on the training data, further complicating the modeling process.

Despite not meeting the expected results initially, this experience has provided valuable insights into the complexities of creating accurate prediction models for medical situations. As we move forward, it will be needed to explore different approaches and utilize the expertise in the field to tackle these challenges and improve the overall performance of the models.

6 Limitation and Suggestion for future study

Despite studying articles and adding new functions for each attribute, the precision and recall rates remained low and were deemed unsatisfactory. However, it soon became clear that it is a difficult task to predict heart strokes using the attribute data in the dataset. Expanding the body data on individuals is required to identify genetic factors, which requires different classifications and datasets for different types of strokes. This includes details on lifestyle elements like diet quality, exercise frequency, and hormone levels like leptin, cortisol, and ghrelin, as well as genetic aspects like MC4R, POMC and cholesterol levels.

The current machine learning recall rates and precision serve as a starting point that needs to be enhanced for future studies. To advance medical research using a variety of datasets, collaboration with medical professionals and doctors is crucial. This project's goal was to give a group of students

a chance to use the knowledge they had learned throughout the semester. Despite the unsatisfactory recall and precision rates, it is hoped that this research will spark interest among others and advance knowledge of heart stroke, ultimately protecting the welfare of our loved ones.

7 Reference

[1] H.-K. Wang et al., “Smoking paradox in stroke survivors? - aha/ASA journals,” Smoking Paradox in Stroke Survivors?, <https://www.ahajournals.org/doi/pdf/10.1161/STROKEAHA.119.027012> (accessed Jun. 30, 2023).

[2] “Diabetes tests,” Centers for Disease Control and Prevention, <https://www.cdc.gov/diabetes/basics/getting-tested.html> (accessed Jun. 30, 2023).

[3] “Stroke,” Stroke | ADA, <https://diabetes.org/diabetes/stroke#:~:text=If%20you%20have%20diabetes%2C%20> (accessed Jun. 30, 2023).

[4] M. Tobias Kurth, “Body mass index and the risk of stroke in men,” Archives of Internal Medicine, <https://jamanetwork.com/journals/jamainternalmedicine/fullarticle/754810> (accessed Jun. 30, 2023).